

Assignment 4: Flower Image Classification Using CNNs

Student Name: Adir Serruya

ID: 316472455

Submission Date: 2/2/2025

Objective

The goal of this assignment is to classify images of flowers into 102 categories using two pre-trained CNNs: **YOLOv5** and **VGG19**. The models were fine-tuned on the Flowers102 dataset.

Dataset

- **Dataset Name:** Flowers102
- **Total Images:** 8,189
- **Classes:** 102
- **Splits:**
 - Training: 50%
 - Validation: 25%
 - Testing: 25%
- **Preprocessing:**
 - Resized images to 224x224.
 - Normalized using ImageNet's mean and standard deviation.

Models

1. YOLOv5

- Adapted for image classification by modifying the final classification layer.
- Fine-tuned with a learning rate of 0.005.

2. VGG19

- Fine-tuned by replacing the last fully connected layer with a 102-class classifier.
- All layers frozen except the last classification layer.

Training Process

1. **Optimizer:** Adam
2. **Loss Function:** Cross-Entropy Loss
3. **Batch Size:** 128
4. **Early Stopping:**
 - Patience of 3 epochs to prevent overfitting.

Results

YOLOv5

- **Validation Accuracy:** 92.97%
- **Test Accuracy:** 93.02%
- **Training Time:** ~28 minutes


VGG19

- **Validation Accuracy:** 76.50%
- **Test Accuracy:** 77.54%
- **Training Time:** ~37 minutes


Plots

1. Loss Over Epochs

YOLOv5




VGG19




2. Accuracy Over Epochs

YOLOv5



VGG19



Conclusion

- **YOLOv5** outperformed **VGG19**, achieving a test accuracy of **93.02%**.
- The results demonstrate the effectiveness of transfer learning for flower classification.
- Further improvements could be made by exploring additional pre-trained models or using ensemble methods.